



PROBABILITY THEORY REVIEW

CSCI-B455



Martha White

DEPARTMENT OF COMPUTER SCIENCE AND INFORMATICS

INDIANA UNIVERSITY, BLOOMINGTON

Spring, 2017

REMINDERS

- Assignment 1 is due on February 1
 - requires reading Chapters 1 and 2
- Thought questions 1 are due on January 25
 - Chapters 1, 2 and 3 (and read the preface)
- Recommendation: do not print out all the notes just yet; later parts will be slightly modified and improved
- See appendix for some background material
 - e.g. a notation sheet
- No class on Monday, January 16 (MLK Jr. day)

PROBABILITY THEORY IS THE SCIENCE OF PREDICTIONS*

- The **goal of science** is to discover theories that can be used to predict how natural processes evolve or explain natural phenomenon, based on observed phenomenon.
- The **goal of probability theory** is to provide the foundation to build theories (= models) that can be used to reason about the outcomes of events, future or past, based on observations.
 - prediction of the unknown which may depend on what is observed and whose nature is probabilistic

*Quote from Csaba Szepesvari, https://eclass.srv.ualberta.ca/pluginfile.php/1136251/mod_resource/content/1/LectureNotes_Probabilities.pdf

(MEASURABLE) SPACE OF OUTCOMES AND EVENTS

Ω = sample space, all outcomes of the experiment

\mathcal{F} = event space, set of subsets of Ω

Ω and \mathcal{F} must be non-empty

If the following conditions hold:

$$1. A \in \mathcal{F} \Rightarrow A^c \in \mathcal{F}$$

$$2. A_1, A_2, \dots \in \mathcal{F} \Rightarrow \bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$$

\mathcal{F} is called a sigma field (sigma algebra)

Note: terminology sigma field sounds technical, but it just means this event space

(Ω, \mathcal{F}) = a measurable space

WHY IS THIS THE DEFINITION?

Intuitively,

1. A collection of outcomes is an event (e.g., either a 1 or 6 was rolled)
2. If we can measure two events separately, then their union should also be a measurable event
3. If we can measure an event, then we should be able to measure that that event did not occur (the complement)

Ω = sample space, all outcomes of the experiment

\mathcal{F} = event space, set of subsets of Ω

If the following conditions hold:

$$1. A \in \mathcal{F} \Rightarrow A^c \in \mathcal{F}$$

$$2. A_1, A_2, \dots \in \mathcal{F} \Rightarrow \bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$$

AXIOMS OF PROBABILITY

(Ω, \mathcal{F}) = a measurable space

Any function $P : \mathcal{F} \rightarrow [0, 1]$ such that

1. (unit measure) $P(\Omega) = 1$
2. (σ -additivity) Any countable sequence of disjoint events $A_1, A_2, \dots \in \mathcal{F}$ satisfies $P(\cup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} P(A_i)$

is called a probability measure (probability distribution)

(Ω, \mathcal{F}, P) = a probability space

WHY NOT THE SIMPLER DEFINITION OF FINITE UNIONS?

In most cases, additivity is enough


$$2. \forall A, B \in \mathcal{F} \text{ and } A \cap B = \emptyset \Rightarrow P(A \cup B) = P(A) + P(B)$$

WHY THESE SEEMINGLY ARBITRARY RULES?

- These rules ensure nice properties of measures
- Other possibilities, these ones chosen

SAMPLE SPACES

Ω



discrete (countable)

continuous (uncountable)

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

$$\Omega = \mathbb{N}$$

$$\text{e.g., } \mathcal{F} = \{\emptyset, \{1, 2\}, \{3, 4, 5, 6\}, \{1, 2, 3, 4, 5, 6\}\}$$

$$\text{Typically: } \mathcal{F} = \mathcal{P}(\Omega)$$



Power set

$$\Omega = [0, 1]$$

$$\Omega = \mathbb{R}$$

$$\text{e.g., } \mathcal{F} = \{\emptyset, [0, 0.5], (0.5, 1.0], [0, 1]\}$$

$$\text{Typically: } \mathcal{F} = \mathcal{B}(\Omega)$$



Borel field

$$\Omega = [0, 1] \cup \{2\} = \text{mixed space}$$

FINDING PROBABILITY DISTRIBUTIONS

(Ω, \mathcal{F}) = a measurable space

Example: $\Omega = \{0, 1\}$
 $\mathcal{F} = \{\emptyset, \{0\}, \{1\}, \Omega\}$

$$P(A) = \begin{cases} 1 - \alpha & A = \{0\} \\ \alpha & A = \{1\} \\ 0 & A = \emptyset \\ 1 & A = \Omega \end{cases} \quad \alpha \in [0, 1]$$

How can we choose P in practice?

Clearly, we cannot do it arbitrarily.

How can we satisfy all constraints?

PROBABILITY MASS FUNCTIONS

Ω = discrete sample space

$$\mathcal{F} = \mathcal{P}(\Omega)$$

Probability mass function:

1. $p : \Omega \rightarrow [0, 1]$
2. $\sum_{\omega \in \Omega} p(\omega) = 1$

The probability of any event $A \in \mathcal{F}$ is defined as


$$P(A) = \sum_{\omega \in A} p(\omega)$$

ARBITRARY PMFs

e.g. PMF for a fair die (table of values)


$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

$$p(\omega) = 1/6 \quad \forall \omega \in \Omega$$



EXERCISE: HOW ARE PMFs USEFUL AS A MODEL?

- Recall we modeled commute times using a gamma distribution (continuous time t)
- Instead could use a probability table for minutes: count number of times $t = 1, 2, 3, \dots$ occurs and then normalize probabilities by # samples
 - why normalize by number of samples?
- Pick t with the largest $p(t)$



USEFUL PMFs

Bernoulli distribution:

$$\Omega = \{S, F\} \quad \alpha \in (0, 1)$$

$$p(\omega) = \begin{cases} \alpha & \omega = S \\ 1 - \alpha & \omega = F \end{cases}$$

Alternatively, $\Omega = \{0, 1\}$

$$p(k) = \alpha^k \cdot (1 - \alpha)^{1-k} \quad \forall k \in \Omega$$

USEFUL PMFs

Binomial distribution:

$$\Omega = \{0, 1, \dots, n\} \quad \alpha \in (0, 1)$$

$$p(k) = \binom{n}{k} \alpha^k (1 - \alpha)^{n-k} \quad \forall k \in \Omega$$

The values are k : the number of successes in a sequence of n independent 0/1 Bernoulli(α) experiments


<http://www.math.uah.edu/stat/apps/BinomialCoinExperiment.html>

USEFUL PMFs

Binomial distribution:

$$\Omega = \{0, 1, \dots, n\} \quad \alpha \in (0, 1)$$

$$p(k) = \binom{n}{k} \alpha^k (1 - \alpha)^{n-k} \quad \forall k \in \Omega$$



USEFUL PMFs


Poisson distribution:

e.g., amount of mail received in a day
number of calls received by call center in an hour

$$p(k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

$$\Omega = \{0, 1, \dots\} \quad \lambda \in (0, \infty)$$

$$\forall k \in \Omega$$



PROBABILITY DENSITY FUNCTIONS

Ω = continuous sample space

$$\mathcal{F} = \mathcal{B}(\Omega)$$

Probability density function:

1. $p : \Omega \rightarrow [0, \infty)$

2. $\int_{\Omega} p(\omega) d\omega = 1$

The probability of any event $A \in \mathcal{F}$ is defined as

$$P(A) = \int_A p(\omega) d\omega.$$

PMFs vs. PDFs

Ω = discrete sample space

Consider a singleton event $\{\omega\} \in \mathcal{F}$, where $\omega \in \Omega$

$$P(\{\omega\}) = p(\omega)$$

Ω = continuous sample space

Consider an interval event $A = [x, x + \Delta x]$, where Δ is small

$$\begin{aligned} P(A) &= \int_x^{x+\Delta x} p(\omega) d\omega \\ &\approx p(x) \Delta x \end{aligned}$$

A FEW COMMENTS ON TERMINOLOGY


- A few new terms, including countable, closure
 - only a small amount of terminology used, can google these terms and learn on your own
 - notations sheet in Appendix of notes
- Countable: integers, rational numbers, ...
- Uncountable: real numbers, intervals, ...
- Why this matters: measures (probability) on these sets is different
- Example: for discrete uniform distribution on $\{0.1, 2.0, 3.6\}$, what is the probability of seeing 3.6?
- Example: for uniform distribution on $[0, 1]$, what is the probability of seeing 0.1?

USEFUL PDFs

Uniform distribution: $\Omega = [a, b]$

$$p(\omega) = \frac{1}{b - a}$$

$$\forall \omega \in [a, b]$$




USEFUL PDFs

Gaussian distribution:

$$\Omega = \mathbb{R} \quad \mu \in \mathbb{R}, \sigma \in \mathbb{R}^+$$

$$p(\omega) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(\omega-\mu)^2} \quad \forall \omega \in \mathbb{R}$$




USEFUL PDFs

Exponential distribution:


$$\Omega = [0, \infty) \quad \lambda > 0$$

$$p(\omega) = \lambda e^{-\lambda\omega}$$


$$\forall \omega \geq 0$$



EXERCISE: MODELING COMMUTE TIMES





Binomial



Poisson

Which might you choose?





Gaussian

$$p(\omega) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(\omega - \mu)^2}$$

EXERCISE: UTILITY OF PDFs AS A MODEL

- Gamma distribution for commute times extrapolates between recorded time in minutes

$$\Gamma(t|k, \theta) = \frac{t^{k-1} e^{-\frac{t}{\theta}}}{\theta^k \Gamma(k)}$$



MULTIDIMENSIONAL PMFs

$$\Omega = \Omega_1 \times \Omega_2 \times \dots \times \Omega_k$$

$$\mathcal{F} = \mathcal{P}(\Omega)$$

Probability mass function:

$$1. \ p : \Omega_1 \times \Omega_2 \times \dots \times \Omega_k \rightarrow [0, 1]$$

$$2. \ \sum_{\omega_1 \in \Omega_1} \cdots \sum_{\omega_k \in \Omega_k} p(\omega_1, \omega_2, \dots, \omega_k) = 1$$

The probability of any event $A \in \mathcal{F}$ is defined as

$$P(A) = \sum_{\omega \in A} p(\omega)$$


$$\omega = (\omega_1, \omega_2, \dots, \omega_k)$$

MULTIDIMENSIONAL PMF

Now record both commute time and number red lights

$$\Omega = \{4, \dots, 14\} \times \{1, 2, 3, 4, 5\}$$

PMF is normalized 2-d table (histogram) of occurrences



MULTIDIMENSIONAL PDFs

$$\Omega = \mathbb{R}^k$$

$$\mathcal{F} = \mathcal{B}(\mathbb{R})^k$$

Probability density function:

1. $p : \mathbb{R}^k \rightarrow [0, \infty)$

2. $\int_{-\infty}^{\infty} \cdots \int_{-\infty}^{\infty} p(\omega_1, \omega_2, \dots, \omega_k) d\omega_1 \cdots d\omega_k = 1$

The probability of any event $A \in \mathcal{F}$ is defined as

$$P(A) = \int_{\omega \in A} p(\omega) d\omega.$$

$\omega = (\omega_1, \omega_2, \dots, \omega_k)$

MULTIDIMENSIONAL GAUSSIAN


$$\Omega = \mathbb{R}^k$$
$$\mathcal{F} = \mathcal{B}(\mathbb{R})^k$$

$$\boldsymbol{\mu} \in \mathbb{R}^k$$

$\boldsymbol{\Sigma}$ = positive definite k -by- k matrix

$|\boldsymbol{\Sigma}|$ = determinant of $\boldsymbol{\Sigma}$

$$p(\boldsymbol{\omega}) = \frac{1}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}} \exp \left(-\frac{1}{2} (\boldsymbol{\omega} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\boldsymbol{\omega} - \boldsymbol{\mu}) \right)$$



QUICK SURVEY

- Who has heard of vectors?
- Who has heard of dot products?
- Who has heard of matrices?

MULTIPLE VARIABLES

- A vector can be thought of as a 1-dimensional array of length d
- A matrix can be thought of as a 2-dimensional array, of dimension $n \times d$

Two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^d$

Dot product $\mathbf{x}^\top \mathbf{y} = \sum_{i=1}^d x_i y_i$

RANDOM VARIABLES

(Ω, \mathcal{F}, P)

Ω

Age: 35
Height: 1.85m
Weight: 75kg
IQ: 104
Likes sports: Yes
Smokes: No
Marital st.: Single
Occupation: Musician

Age: 26
Height: 1.75m
Weight: 79kg
IQ: 103
Likes sports: Yes
Smokes: No
Marital st.: Divorced
Occupation: Athlete

$$A = \{\omega \in \Omega : \text{Musician}(\omega) = \text{yes}\}$$

Musician is a random variable (a function)

A is the new event space

Can ask $P(M = 0)$ and $P(M = 1)$

WE INSTINCTIVELY CREATE THIS TRANSFORMATION

Assume Ω is a set of people.

Compute the probability that a randomly selected person $\omega \in \Omega$ has a cold.

Define event $A = \{\omega \in \Omega : \text{Disease}(\omega) = \text{cold}\}$.

Disease is our new random variable, $P(\text{Disease} = \text{cold})$

Disease is a function that maps outcome space to new outcome space $\{\text{cold}, \text{not cold}\}$

RANDOM VARIABLES



Example: three consecutive (fair) coin tosses
 X = the number of heads in the first toss
 Y = the number of heads in all three tosses
Find the probability spaces after the transformations.

Where is the probability space (Ω, \mathcal{F}, P) ?

Where is the randomness?

$$\Omega = \{HHH, HHT, HTH, HTT, THH, THT, TTH, TTT\}$$

$$\mathcal{F} = \mathcal{P}(\Omega)$$

$$P = ?$$

$$P(\Omega) = 1$$

$$P(\{HHH, TTT\}) = \frac{2}{8}$$

$$\vdots$$

RANDOM VARIABLES



$$X : \Omega \rightarrow \{0, 1\}$$

$$Y : \Omega \rightarrow \{0, 1, 2, 3\}$$

ω	HHH	HHT	HTH	HTT	THH	THT	TTH	TTT
$X(\omega)$	1	1	1	1	0	0	0	0
$Y(\omega)$	3	2	2	1	2	1	1	0

What are the probability spaces $(\Omega_X, \mathcal{F}_X, P_X)$ and $(\Omega_Y, \mathcal{F}_Y, P_Y)$?

Where does the randomness come from?

Once we have these new spaces, same pdf and pdf definitions apply

RANDOM VARIABLE: FORMAL DEFINITION

(Ω, \mathcal{F}, P) = a probability space

Random variable:

1. $X : \Omega \rightarrow \Omega_X$

2. $\forall A \in \mathcal{B}(\Omega_X)$ it holds that $\{\omega : X(\omega) \in A\} \in \mathcal{F}$

It follows that: $P_X(A) = P(\{\omega : X(\omega) \in A\})$

Example $X : \Omega \rightarrow [0, \infty)$

Ω is set of (measured) people in population

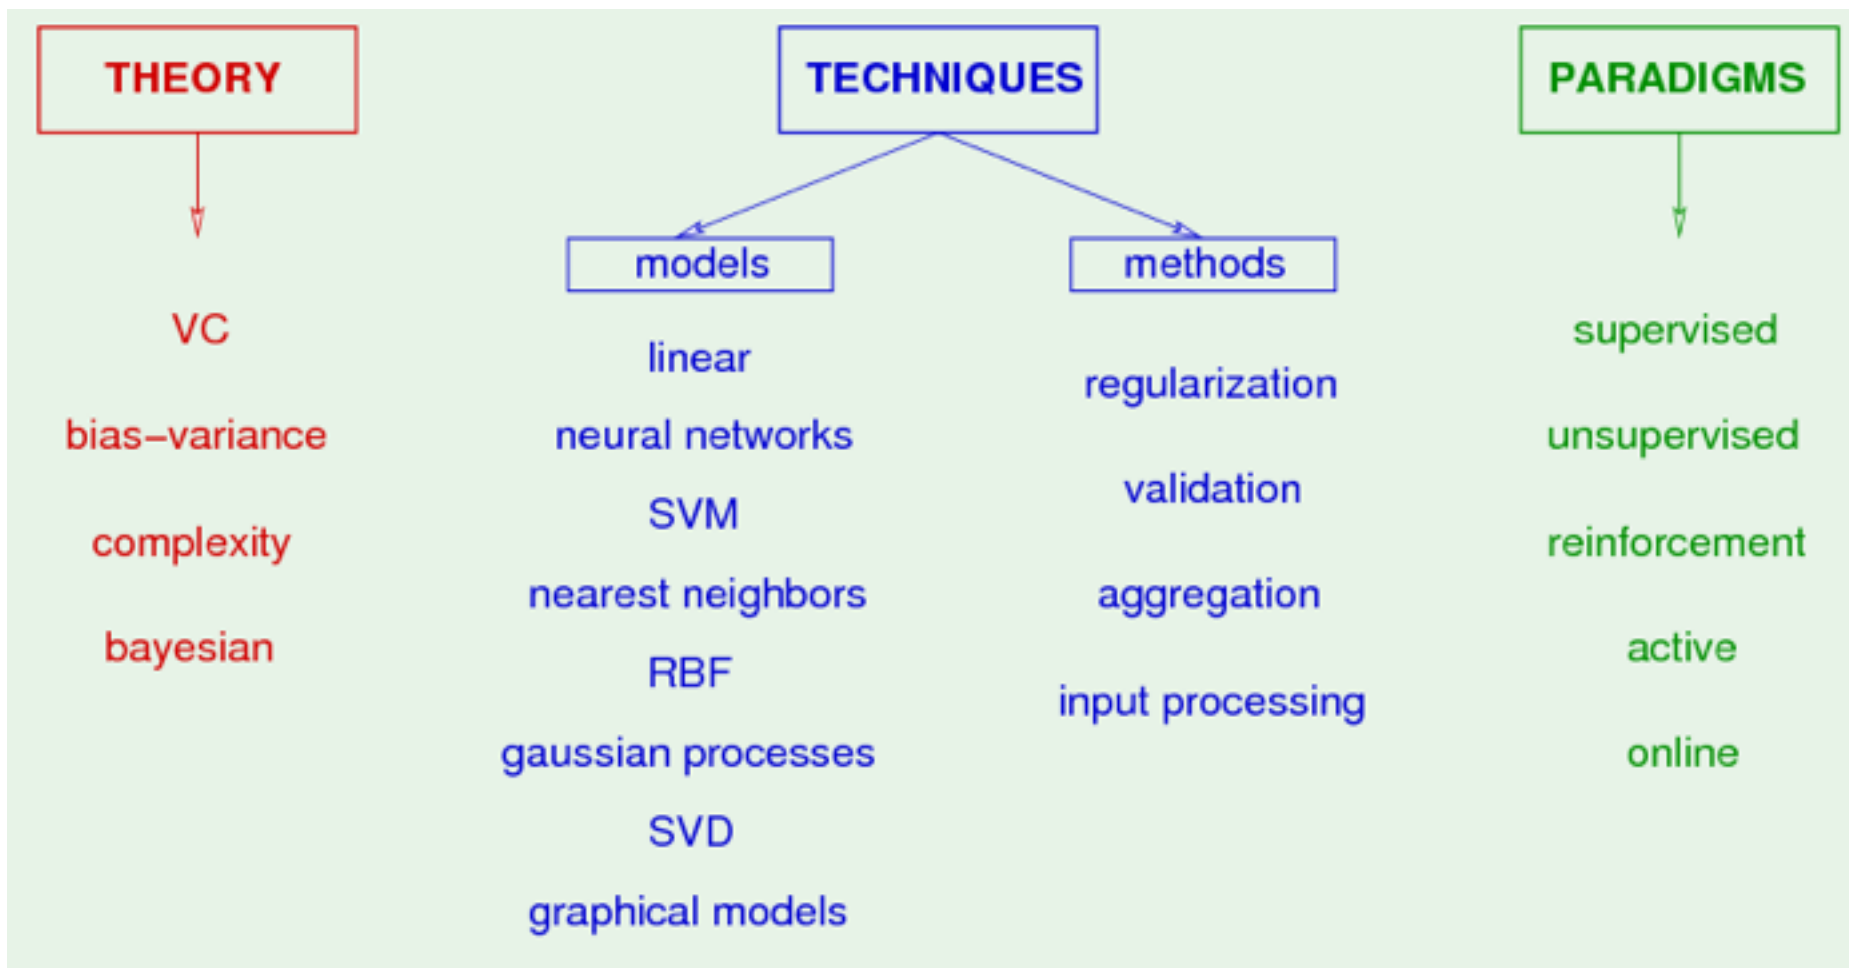
with associated measurements such as height and weight

$X(\omega)$ = height

A = interval = $[5'1'', 5'2'']$

$$P(X \in A) = P(5'1'' \leq X \leq 5'2'') = P(\{\omega : X(\omega) \in A\})$$

JAN. 18: PROBABILITY REVIEW CONTINUED



Machine learning topic overview

* from Yaser Abu-Mostafa, <https://work.caltech.edu/library/>

REMINDERS

- Assignment 1 is due on February 1
- Thought questions 1 are due on January 25
- Office hours
 - Martha: 3-5 p.m. on Tuesday (LH 401E)
 - Inhak: 3:30 - 5:30 p.m. on Tuesday (LH 325)
 - Andrew: 12:00 - 2:00 p.m. on Thur (LH 215D)
- Up-front background material
 - Immersion style: understanding more in-depth as we use the ideas from probability
- Lecture notes posted before class
- I do not expect you to know formulas, like pdfs

KEY POINTS SO FAR

- Many of our variables will be random
- These random variables can be discrete or continuous
 - discrete e.g. $\{0, 1, 2\}$
 - continuous, e.g. $[-100, 100]$
- Several named PMFs and PDFs to provide distributions over the possible values
 - why? explicit functional form will be useful later
 - e.g., $p(x) = \lambda \exp(-\lambda x)$
- Multi-variate distributions natural extensions of scalar distributions; probabilities over vector instances, e.g., x in $[-10, 10]^2$

CONDITIONAL DISTRIBUTIONS

Conditional probability distribution:

$$p_{Y|X}(y|x) = \frac{p_{XY}(x, y)}{p_X(x)}$$

The probability of an event A , given that $X = x$, is:

$$P_{Y|X}(Y \in A|X = x) = \begin{cases} \sum_{y \in A} p_{Y|X}(y|x) & Y : \text{discrete} \\ \int_{y \in A} p_{Y|X}(y|x) dy & Y : \text{continuous} \end{cases}$$

DROPPING SUBSCRIPTS

Instead of:

$$p_{Y|X}(y|x) = \frac{p_{XY}(x, y)}{p_X(x)}$$

We will write:

$$p(y|x) = \frac{p(x, y)}{p(x)}$$

EXAMPLE


- Let X be a Bernoulli random variable (i.e., 0 or 1 with probability α)
- Let Y be a random variable in $\{10, 11, \dots, 1000\}$
- $p(y \mid X = 0)$ and $p(y \mid X = 1)$ are different distributions
- Two types of books: fiction ($X=0$) and non-fiction ($X=1$)
- Let Y corresponds to number of pages
- Distribution over number of pages different for fiction and non-fiction books (e.g., average different)

EXAMPLE CONTINUED

- Two types of books: fiction ($X=0$) and non-fiction ($X=1$)
- Y corresponds to number of pages
- $p(y \mid X = 0) = p(X = 0, y)/p(X = 0)$
- $p(X = 0, y)$ = probability that a book is fiction and has y pages (imagine randomly sampling a book)
- $p(X = 0)$ = probability that a book is fiction
- If most books are non-fiction, $p(X = 0, y)$ is small even if y is a likely number of pages for a fiction book
- $p(X = 0)$ accounts for the fact that joint probability small if $p(X = 0)$ is small

ANOTHER EXAMPLE

- Two types of books: fiction ($X=0$) and non-fiction ($X=1$)
- Let Y be a random variable over the reals, which corresponds to amount of money made
- $p(y \mid X = 0)$ and $p(y \mid X = 1)$ are different distributions
- e.g., even if both $p(y \mid X = 0)$ and $p(y \mid X = 1)$ are Gaussian, they likely have different means and variances



WHAT DO WE KNOW ABOUT $P(Y)$?

- We know $p(y \mid x)$
- We know marginal $p(x)$
- Correspondingly we know $p(x, y) = p(y \mid x) p(x)$
 - from conditional probability definition that $p(y \mid x) = p(x, y) / p(x)$
- What is the marginal $p(y)$?

$$\begin{aligned} p(y) &= \sum_x p(x, y) \\ &= \sum_x p(y|x)p(x) \\ &= p(y|X = 0)p(X = 0) + p(y|X = 1)p(X = 1) \end{aligned}$$

CHAIN RULE

Conditional probability distribution:

$$p(x_k | x_1, \dots, x_{k-1}) = \frac{p(x_1, \dots, x_k)}{p(x_1, \dots, x_{k-1})}$$


This leads to:

$$p(x_1, \dots, x_k) = p(x_1) \prod_{l=2}^k p(x_l | x_1, \dots, x_{l-1})$$

Two variable example $p(x, y) = p(x|y)p(y) = p(y|x)p(x)$

EXERCISE: CONDITIONAL PROBABILITIES

- Using conditional probabilities, we can incorporate other external information (features)
- Let y be the commute time, x the day of the year
- Array of conditional probability values $\rightarrow p(y \mid x)$
 - $y = 1, 2, \dots$ and $x = 1, 2, \dots, 365$
- What are some issues with this choice for x ?
- What other x could we use feasibly?




EXERCISE: ADDING IN AUXILIARY INFORMATION

- Gamma distribution for commute times extrapolates between recorded time in minutes
- Can incorporate external information (features) by modeling $\theta = \text{function}(\text{features})$

$$\theta = \sum_{i=1}^d w_i x_i$$

$$\Gamma(t|k, \theta) = \frac{t^{k-1} e^{-\frac{t}{\theta}}}{\theta^k \Gamma(k)}$$



INDEPENDENCE OF RANDOM VARIABLES

X and Y are **independent** if:

$$p(x, y) = p(x)p(y)$$

X and Y are **conditionally independent** given Z if:

$$p(x, y|z) = p(x|z)p(y|z)$$


Exercise: what if we had k random variables, X_1, \dots, X_k ?

CONDITIONAL INDEPENDENCE EXAMPLES

- Imagine you have a biased coin (does not flip 50% heads and 50% tails, but skewed towards one)
- Let Z = bias of a coin (say outcomes are 0.3, 0.5, 0.8 with associated probabilities 0.7, 0.2, 0.1)
 - what other outcome space could we consider?
 - what kinds of distributions?
- Let X and Y be consecutive flips of the coin
- Are X and Y independent?
- Are X and Y conditionally independent, given Z ?

EXPECTED VALUE (MEAN)


$$\mathbb{E}[X] = \begin{cases} \sum_{x \in \mathcal{X}} xp(x) & X : \text{discrete} \\ \int_{\mathcal{X}} xp(x)dx & X : \text{continuous} \end{cases}$$



EXPECTED VALUE FOR MULTIVARIATE

$$\mathbb{E} [\mathbf{X}] = \begin{cases} \sum_{\mathbf{x} \in \mathcal{X}} \mathbf{x} p(\mathbf{x}) & \mathbf{X} : \text{discrete} \\ \int_{\mathcal{X}} \mathbf{x} p(\mathbf{x}) d\mathbf{x} & \mathbf{X} : \text{continuous} \end{cases}$$

Each instance \mathbf{x} is a vector, p is a function on these vectors



CONDITIONAL EXPECTATIONS

$$\mathbb{E}[Y|X = x] = \begin{cases} \sum_{y \in \mathcal{Y}} yp(y|x) & Y : \text{discrete} \\ \int_{\mathcal{Y}} yp(y|x)dy & Y : \text{continuous} \end{cases}$$


Different expected value, depending on which x is observed

EXERCISE: RVs, PDFs AND UNCERTAINTY


- In ML, common strategy to assume trying to learn a deterministic function, from noisy measurements
- Denoised “truth”: $y = f(x)$
- Noisy observation: $f(x) + \text{noise}$
 - one common assumption is the noise N is a Gaussian RV
 - $E[f(x) + \text{noise}] = f(x) + E[\text{noise}] = f(x) = E[Y \mid x]$
- For a sample x of RV X :

$$N \sim \mathcal{N}(0, \sigma^2)$$

$$Y = f(x) + N \sim \mathcal{N}(f(x), \sigma^2)$$



COVARIANCE



$$\begin{aligned}\text{Cov}[X, Y] &= \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])] \\ &= \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y],\end{aligned}$$

$$\text{Corr}[X, Y] = \frac{\text{Cov}[X, Y]}{\sqrt{V[X] \cdot V[Y]}},$$

COVARIANCE FOR MORE THAN TWO DIMENSIONS

$$\mathbf{X} = [X_1, \dots, X_d]$$

$$\begin{aligned}\Sigma_{ij} &= \text{Cov}[X_i, X_j] \\ &= \mathbb{E}[(X_i - \mathbb{E}[X_i])(X_j - \mathbb{E}[X_j])]\end{aligned}$$

$$\begin{aligned}\mathbf{\Sigma} &= \text{Cov}[\mathbf{X}, \mathbf{X}] \\ &= \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^\top] \\ &= \mathbb{E}[\mathbf{X}\mathbf{X}^\top] - \mathbb{E}[\mathbf{X}]\mathbb{E}[\mathbf{X}]^\top.\end{aligned}$$

COVARIANCE FOR MORE THAN TWO DIMENSIONS

$$\mathbf{X} = [X_1, \dots, X_d]$$

$$\begin{aligned}\boldsymbol{\Sigma} &= \text{Cov}[\mathbf{X}, \mathbf{X}] \\ &= \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^\top] \\ &= \mathbb{E}[\mathbf{X}\mathbf{X}^\top] - \mathbb{E}[\mathbf{X}]\mathbb{E}[\mathbf{X}]^\top.\end{aligned}$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d$$

Dot product

$$\mathbf{x}^\top \mathbf{y} = \sum_{i=1}^d x_i y_i$$

Outer product

$$\mathbf{xy}^\top = \begin{bmatrix} x_1 y_1 & x_1 y_2 & \dots & x_1 y_d \\ x_2 y_1 & x_2 y_2 & \dots & x_2 y_d \\ \vdots & \vdots & & \vdots \\ x_d y_1 & x_d y_2 & \dots & x_d y_d \end{bmatrix}$$

SOME USEFUL PROPERTIES

1. $\mathbb{E}[c\mathbf{X}] = c\mathbb{E}[\mathbf{X}]$
2. $\mathbb{E}[\mathbf{X} + \mathbf{Y}] = \mathbb{E}[\mathbf{X}] + \mathbb{E}[\mathbf{Y}]$
3. $V[c] = 0$ \triangleright the variance of a constant is zero
4. $V[\mathbf{X}] \succeq 0$ (i.e., is positive semi-definite), where for $d = 1$, $V[\mathbf{X}] \geq 0$
 $V[\mathbf{X}]$ is shorthand for $\text{Cov}[\mathbf{X}, \mathbf{X}]$.
5. $V[c\mathbf{X}] = c^2 V[\mathbf{X}]$.
6. $\text{Cov}[\mathbf{X}, \mathbf{Y}] = \mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{Y} - \mathbb{E}[\mathbf{Y}])^\top] = \mathbb{E}[\mathbf{X}\mathbf{Y}^\top] - \mathbb{E}[\mathbf{X}]\mathbb{E}[\mathbf{Y}]^\top$
7. $\text{Cov}[\mathbf{X} + \mathbf{Y}] = V[\mathbf{X}] + V[\mathbf{Y}] + 2\text{Cov}[\mathbf{X}, \mathbf{Y}]$

EXAMPLE: SAMPLE AVERAGE IS UNBIASED ESTIMATOR

Obtain instances x_1, \dots, x_n

What can we say about the sample average?

This sample is random, so we consider i.i.d. random variables X_1, \dots, X_n

Reflects that we could have seen a different set of instances x_i

$$\begin{aligned}\mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n X_i \right] &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] \\ &= \frac{1}{n} \sum_{i=1}^n \mu \\ &= \mu\end{aligned}$$

For any one sample x_1, \dots, x_n , unlikely that $\frac{1}{n} \sum_{i=1}^n x_i = \mu$

MIXTURES OF DISTRIBUTIONS

Mixture model:

A set of m probability distributions, $\{p_i(x)\}_{i=1}^m$

$$p(x) = \sum_{i=1}^m w_i p_i(x)$$

where $\mathbf{w} = (w_1, w_2, \dots, w_m)$ and non-negative and


$$\sum_{i=1}^m w_i = 1$$

MIXTURES OF GAUSSIANS

Mixture of $m = 2$ Gaussian distributions:

$$w_1 = 0.75, w_2 = 0.25$$

$$p(x) = \sum_{i=1}^m w_i p_i(x)$$



SUMMARY: PARAMETRIC MODELS

- We will consider many parametric models in machine learning
- To model the data, we can pick a parametric class and do parameter estimation (next)
- Given a model, we can make statements about our data
 - predict target given inputs (conditional probs)
 - find underlying structure of data
 - find explanatory variables
 - ...
- We will incrementally generalize the types of models we consider to model our data