



COMPUTER SCIENCE

INDIANA UNIVERSITY


School of Informatics and Computing
Bloomington

Performance measures



Crash course in evaluation

- Running a scientific experiment to compare machine learning algorithms necessary to draw conclusions
- Needs to be meticulous, even if sometimes tedious or need to re-run experiments
- Requires an experiment design and statistical significance tests





Experimental set-up

- Performance measures
- Sampling: How to obtain multiple samples of performance?
- Making conclusions: Statistical significance tests
- Careful statistical work done on executing empirical studies; pros and cons to each
 - for a nice reference, see Evaluating Learning Algorithms: A Classification Perspective (http://www.mohakshah.com/tutorials/icml2012/Tutorial-ICML2012/Tutorial_at_ICML_2012.html); slides in this lecture use some of the material there
 - “Prediction error estimation: a comparison of resampling methods”



Statistical significant test

- Can the observed results be attributed to real characteristics of the learner under scrutiny or are they observed by chance?
- Hypothesis testing:
 - State a null hypothesis, e.g., the expected errors of two classifiers is equivalent
 - Choose a statistical significance test to reject the null hypothesis; failing to reject the null hypothesis does not mean we accept it
 - Rejecting the null hypothesis gives us some confidence in the belief that our observations did not occur merely by chance.

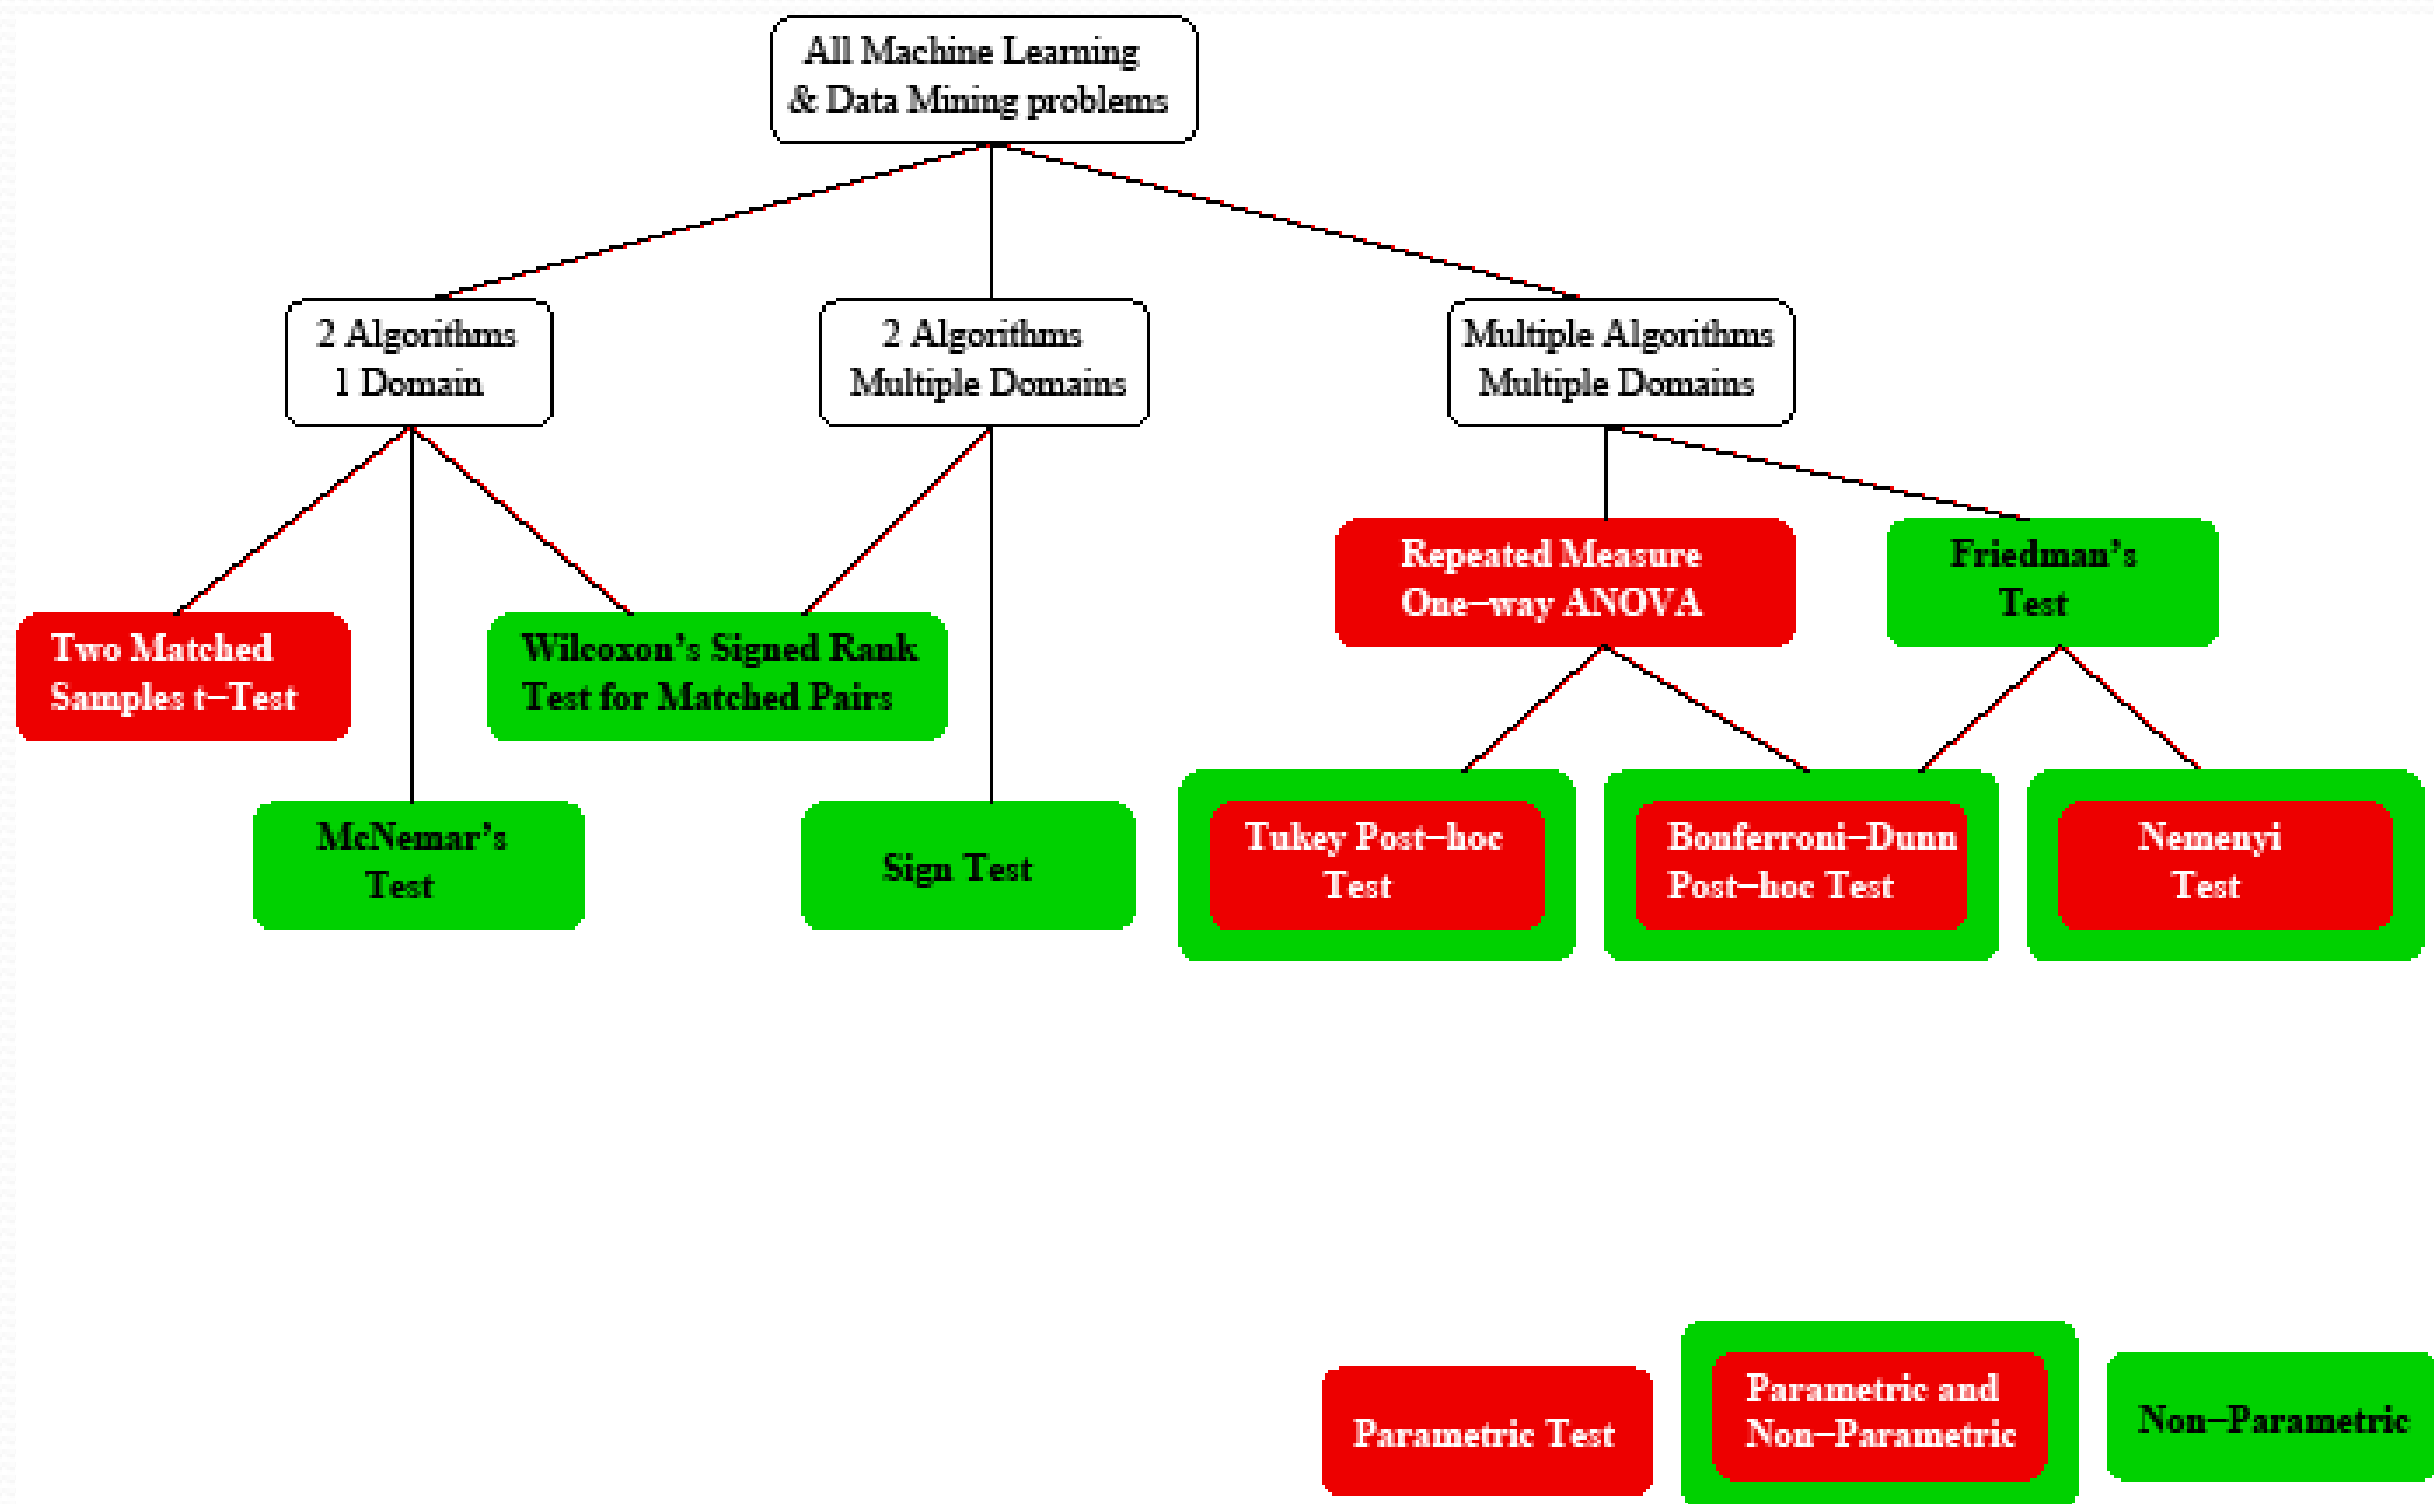


How to choose tests?

- Try to satisfy assumptions and use some rules of thumb
- Parametric statistical tests make stronger assumptions about the distribution of the data
- Non-parametric tests make weaker assumptions, but are less powerful (less able to reject the null hypothesis when it is false)
- Selection based on type of problem
 - comparing 2 algorithms on a single domain
 - comparing 2 algorithms across domains
 - comparing multiple algorithms across domains



Statistical test summary







Comparing two algorithms on a single domain

- Imagine you have N independent test-samples, giving N paired measures of error for the two algorithms
- Simplest (not very powerful) strategy:
 - compute two (95%) confidence intervals for the means
 - if the two intervals do not overlap, means are significantly different

different




unknown





More powerful strategy: t-test

- Confidence intervals may overlap, but the means may still be statistically different
- Paired t-test enables a more powerful comparison
 - more ability to reject the null hypothesis






Assumptions of the t-test

- **The Normality or Pseudo-Normality Assumption:** samples come from normally distributed population. Alternatively, the sample size of the testing set should be greater than 30.
- **The Randomness of the Samples:** The sample should be representative of the underlying population. Therefore, the instances of the testing set should be randomly chosen from their underlying distribution.
- **Equal Variance of the populations:** The two samples come from populations with equal variance.



Example where assumptions of t-test violated

- **Equal Variance:** variance of C4.5 and NB cannot be considered equal.
- Not warranted to use the t-test to compare C4.5 to NB on the Labour data.
- A better test to use is the non-parametric alternative to the t-test: McNemar's Test



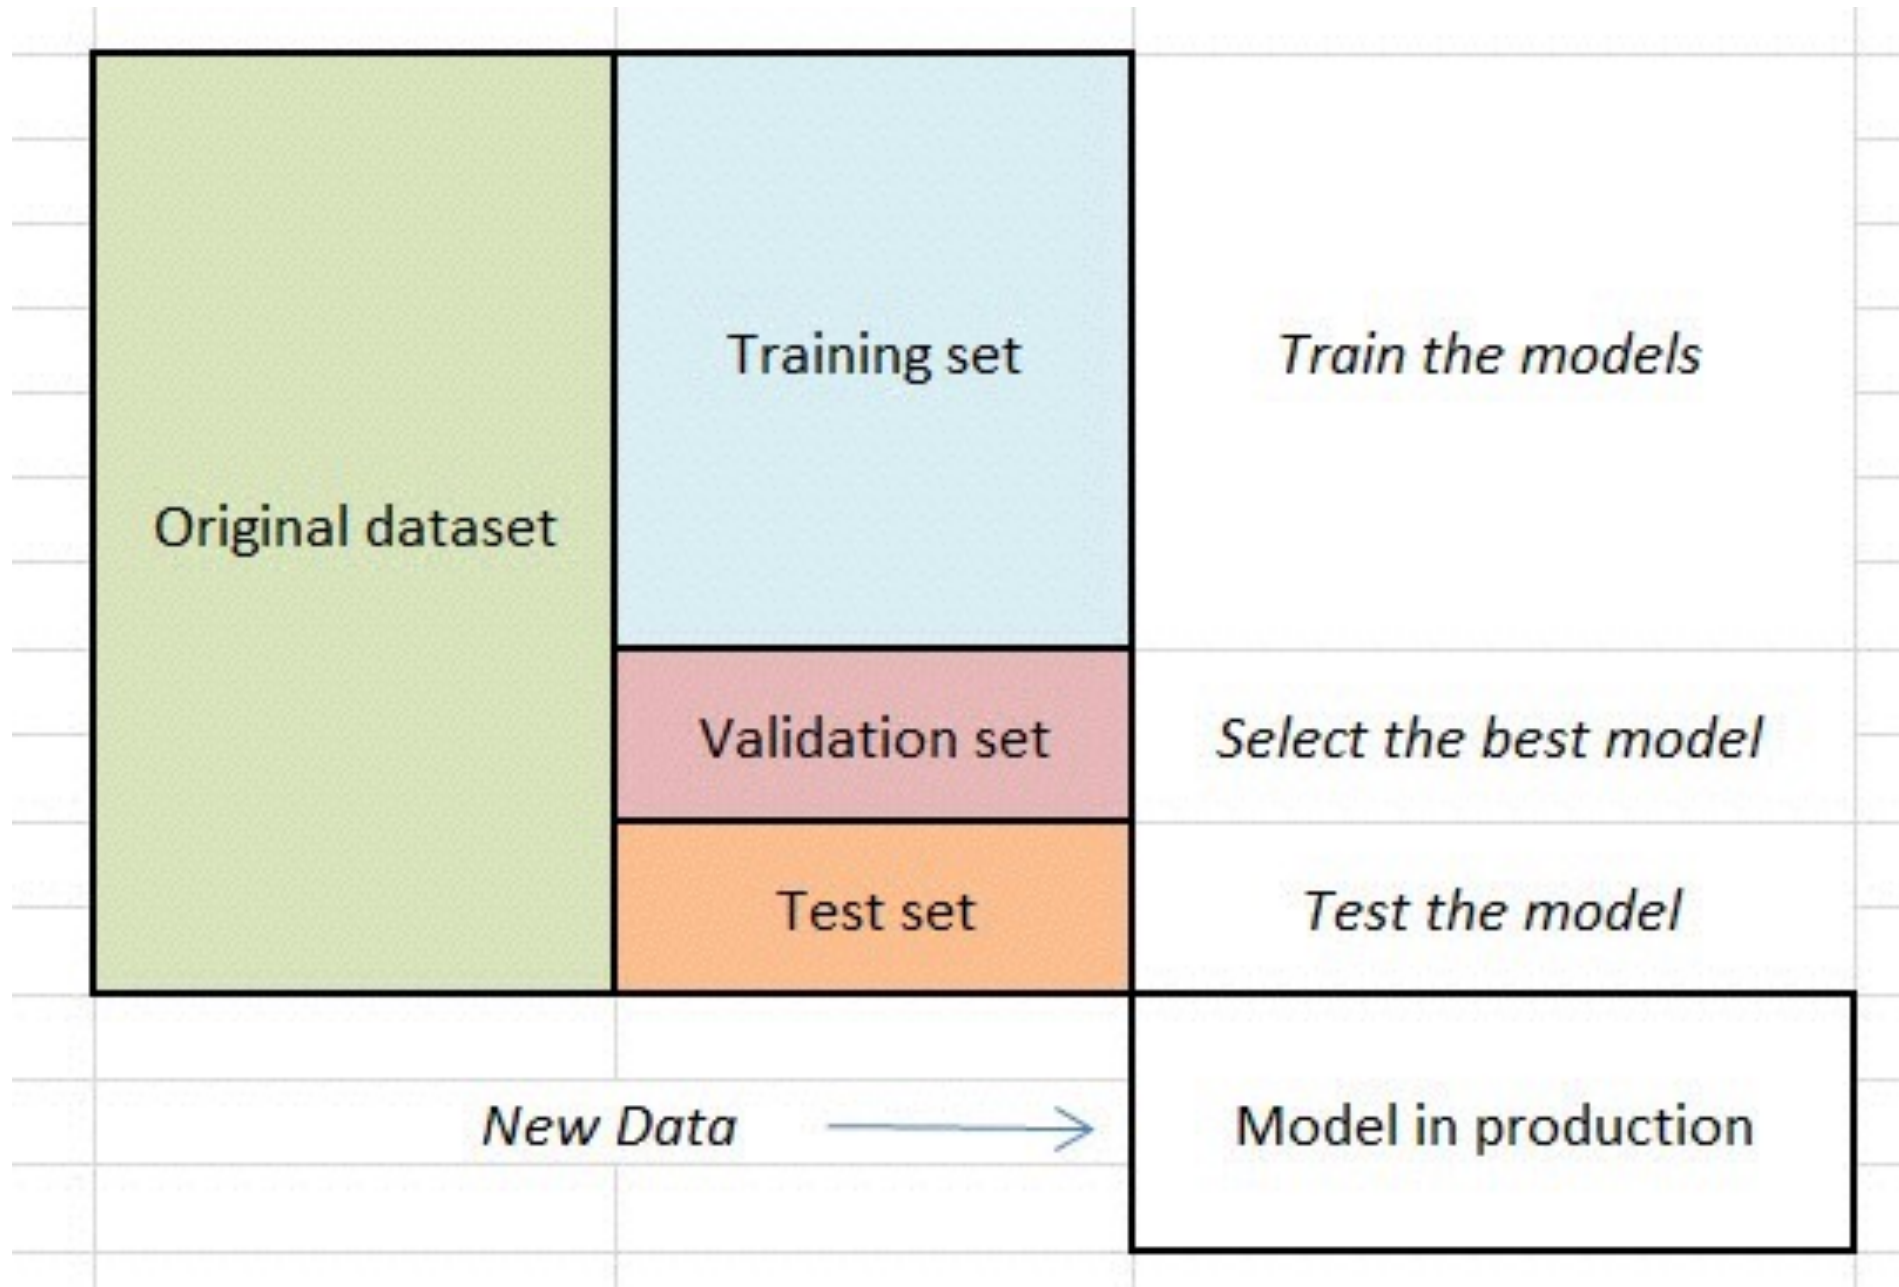


Sampling

- How do we get independent samples of test error? And why?
 - want to get a measure of generalization performance for an algorithm
 - want to compare algorithms
 - want to do hyperparameter selection (e.g., regularization parameters)
- Hold-out test set
- Subsampling approaches
 - Cross-validation (CV)
 - Repeated subsampling: Monte carlo CV
 - Bootstrap resampling



Train-Validation-Test





With lots of data...

- Can afford to have a large hold-out test set that is never used for training
- Data split into train and split
 - Training set may be further split into train and validation, to provide a means to select hyperparameters
- Even with only one split, with a large amount of data likely to get a reasonable estimate of generalization error
- But this approach is error prone; can be difficult to ensure that did not accidentally peak at the data



Thought exercise

- Imagine you learn on the training set for a fixed hyper parameter setting, and get an error measure on the validation set. Do you expect this to overestimate or underestimate your generalization error?
- Imagine you
 - pick a regularization weight using the validation set
 - then you check your test error and notice it is low
 - then you expand your range of regularization weights to improve performance
 - Do you expect the error on the validation set to be an overestimate or underestimate of the generalization error?
 - What about the test error?



Monte carlo CV

- Also called “repeated learning testing-model”
- Randomly sample without replacement the training set and the test set
 - or for smaller datasets, first sample the training set and use the rest for test
- Repeat this random subsample m times to obtain m training/test splits





k-fold CV

Randomly and evenly split into 4 non-overlapping partitions


D

20 data points




Partition 1.

Data points: 1, 3, 5, 15, 16




Partition 2.

Data points: 6, 10, 11, 14, 17



Partition 3.

Data points: 4, 9, 12, 19, 20



Partition 4.

Data points: 2, 7, 8, 13, 17

- Learn model on $k-1$ folds and test on the hold-out fold (done k times); average k error estimates



Internal and External CV

- Internal CV used to select hyperparameters
 - this can be thought of as part of the algorithm, not actually used to evaluate the differences
 - let data decide what hyperparameters to use, so need reasonable subsampling approach to evaluate selection
- External CV used to obtain performance measures
 - e.g., properly compare two algorithms



Example

- For training set i , test hyperparameters using (internal) k -fold CV
- Select the “best” hyperparameters according to average error
- Given these “best” parameters, learn on entire training set i
- Obtain test error on test set i
- This gives you an estimate of error for training/test split i
- Over all m splits, get m estimates of error (different from k)



Which sampling approach should I use?

- No definitive answer, mostly empirical support
- For how to select k , bias-variance trade-off
 - for small k , high-bias and low-variance
 - for large k , low bias but high variance (e.g., leave-one-out)
 - Some experiments showing that a reasonable balance is $k = 10$
 - Also determined by computational resources; large k expensive
- For how to select between sampling methods,
 - repeated CV and Monte Carlo CV shown to have fewer Type 1 errors
- Criteria for internal and external CV may be quite different



Internal

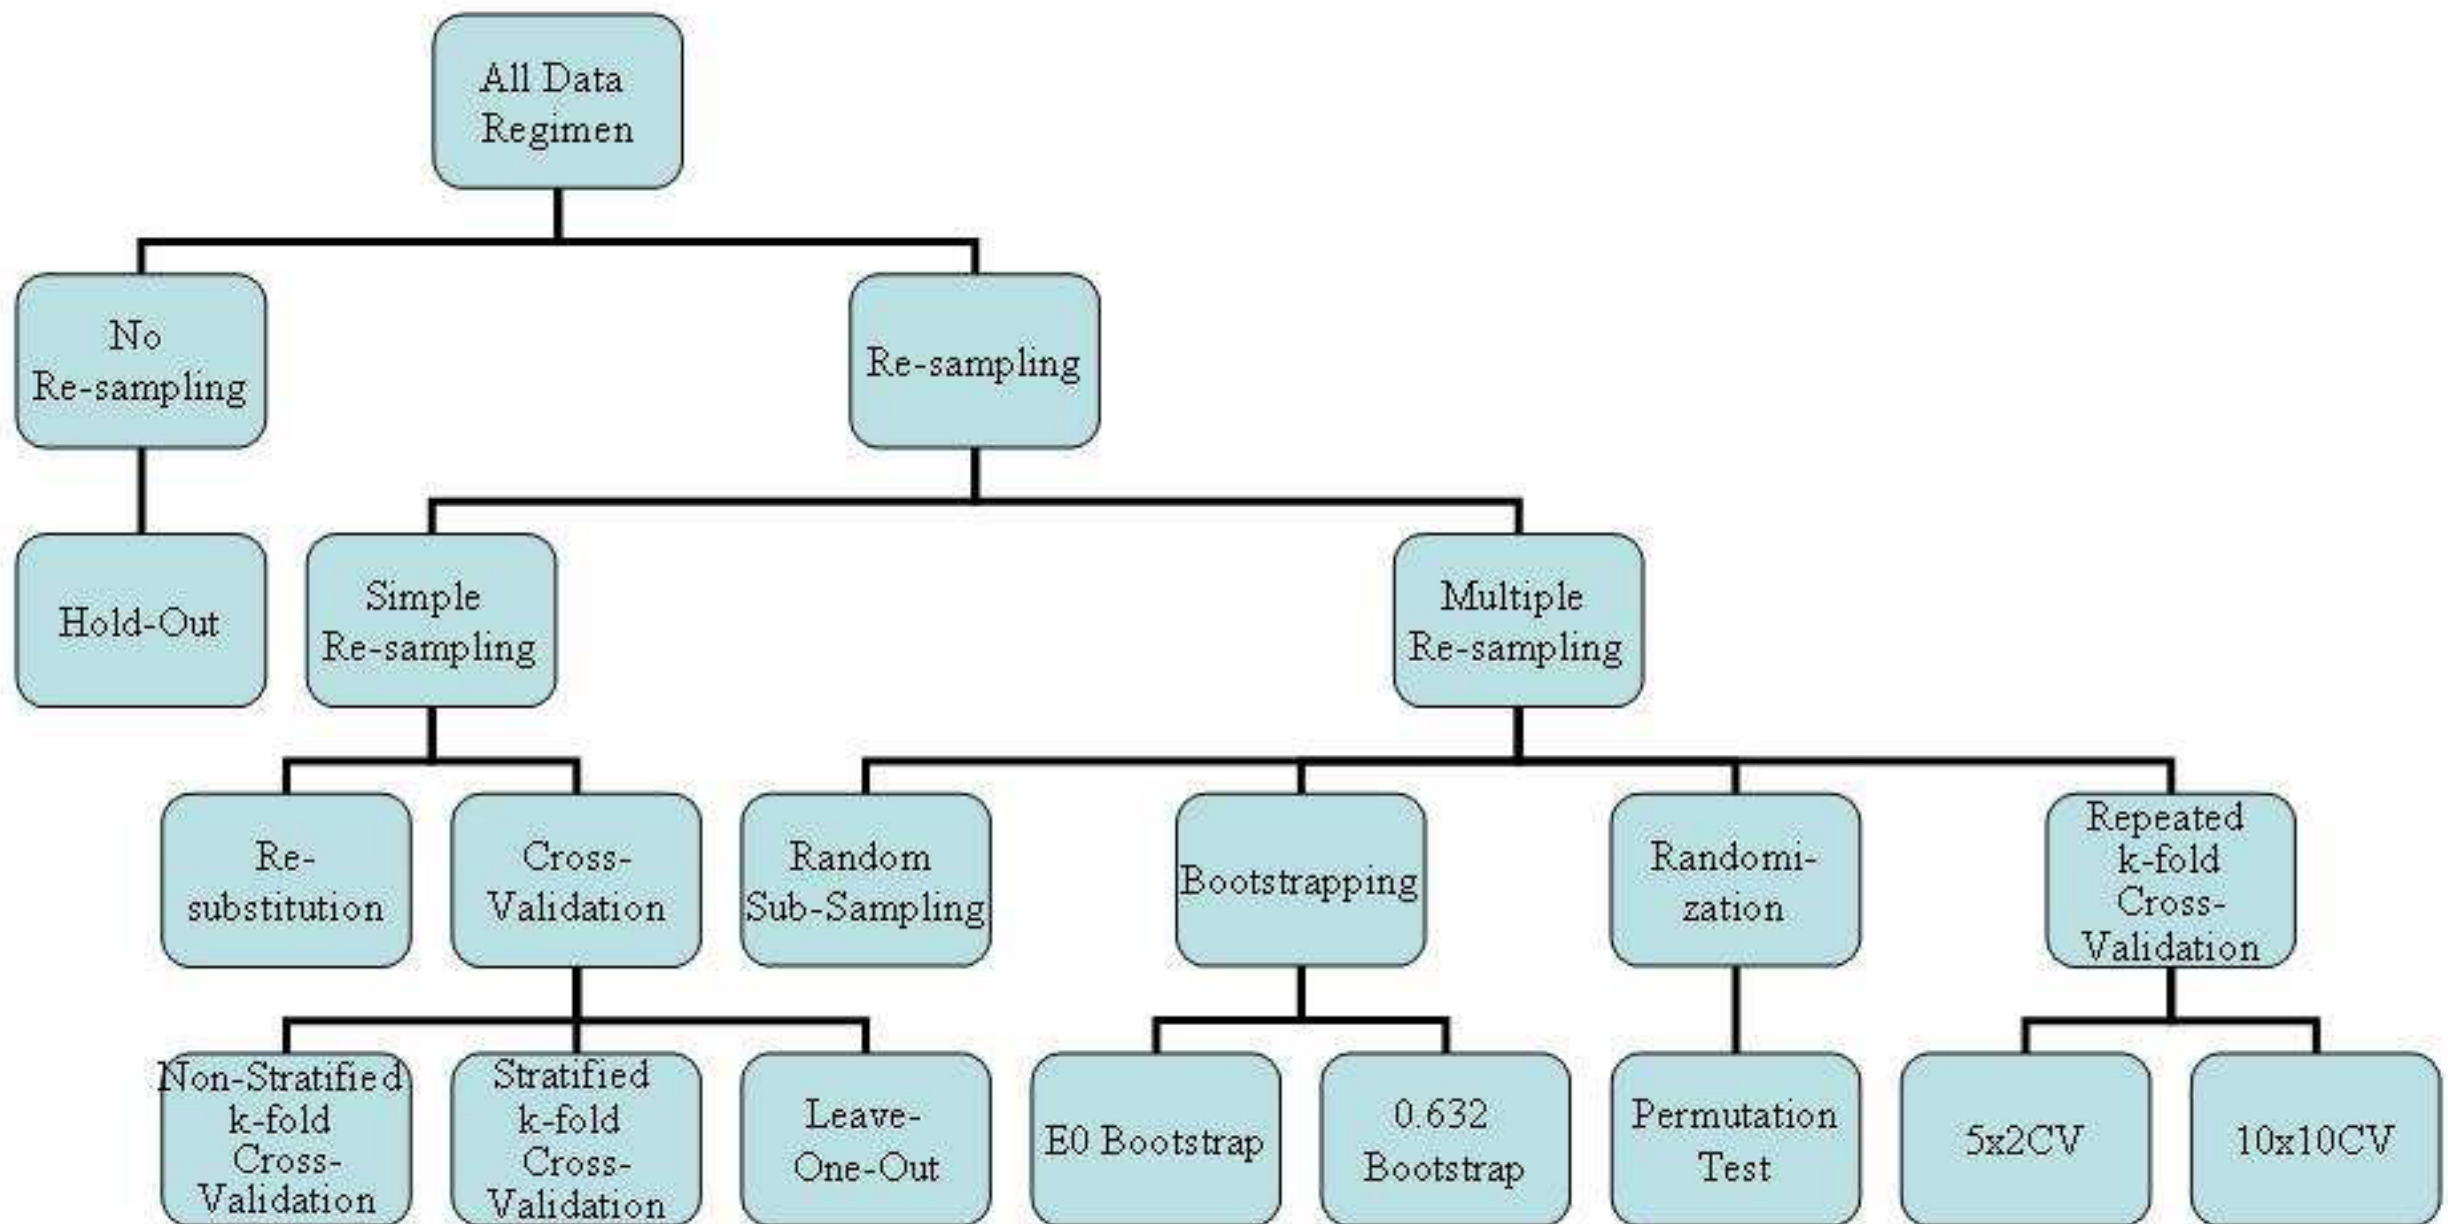
- Training k models can be expensive; want smaller k
- k -fold CV a reasonable choice because gives an almost unbiased estimate of accuracy

External

- Want to use hypothesis testing, e.g., Null Hypothesis is that the means of these two algorithms is the same
- Want a sampling technique that has less Type 1 errors
- Assumptions require independent samples of error, but empirically k -fold is not necessarily better than repeated sampling



Re-sampling summary






Regression objectives

- We have looked at l_2 error for estimating parameters (i.e., as an objective) and to measure performance
- Other options:
 - l_1 error — can be difficult to optimize, still a useful measure of error
 - smooth l_1 — smooth and convex, easier to optimize, not usually used as a measure of error (unless reporting accuracy of optimizer)
 - R-squared — coefficient of determination
 - Variance unexplained
 - Percentage error — rescale by magnitude of values



R-squared measure

- Also called “coefficient of determination” $R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$



- The sum of squares of residuals, also called the **residual sum of squares**:

$$SS_{\text{res}} = \sum_i (y_i - f_i)^2$$

- The **total sum of squares** (proportional to the **variance** of the data):

$$SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2,$$

Larger R-squared is better



R-squared is monotone in number of features

- As add more features, the R-squared measure cannot decrease. Why?

$$R^2 \equiv 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}, \quad SS_{\text{res}} = \sum_i (y_i - f_i)^2 \quad SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2,$$

- Is this an issue?
- Alternative: adjusted R-squared — penalize the number of explanatory variables (features)



Percentage error

- If use error $\| \text{val1} - \text{val2} \|$, and get 0.1, is this good?
- One option: mean percentage error (issues?)
- Another option: mean absolute percentage error (MAPE)

$$M = \frac{1}{n} \sum_{t=1}^n \left| \frac{A_t - F_t}{A_t} \right|,$$

- Another option: symmetric MAPE

$$\text{SMAPE} = \frac{1}{n} \sum_{t=1}^n \frac{|F_t - A_t|}{(|A_t| + |F_t|)/2}$$



Classification terminology

- True positives — samples predicted by classifier to be positive that have true label positive
- False positives — samples predicted by classifier to be positive that have true label negative
- True negatives — samples predicted by classifier to be negative that have true label negative
- False negatives — samples predicted by classifier to be negative that have true label positive



Classification measures

Name	Symbol	Definition
Classification error	$error$	$error = \frac{fp+fn}{tp+fp+tn+fn}$
Classification accuracy	$accuracy$	$accuracy = 1 - error$
True positive rate	tpr	$tpr = \frac{tp}{tp+fn}$
False negative rate	fnr	$fnr = \frac{fn}{tp+fn}$
True negative rate	tnr	$tnr = \frac{tn}{tn+fp}$
False positive rate	fpr	$fpr = \frac{fp}{tn+fp}$
Precision	pr	$pr = \frac{tp}{tp+fp}$
Recall	rc	$rc = \frac{tp}{tp+fn}$



Why these specific values?

- These measures exist for multiple reasons
- Separate the importance of false positives and false negatives
 - In some cases, much more hazardous to have a false positive than a false negative (or vice versa)
- Avoid issues with imbalanced datasets



Confusion Matrix for binary classification

Predicted class	True class	
	0	1
0	N_{00} 😊	N_{01} 😞
1	N_{10} 😞	N_{11} 😊

$$Accuracy = \frac{N_{00} + N_{11}}{N_{00} + N_{10} + N_{01} + N_{11}}$$

Number of data points whose true class was 0 but predicted class was 1.


$$Error = 1 - Accuracy$$



Example of importance of measures: imbalanced datasets

16 data points have class 0 (majority class)

4 data points have class 1 (minority class)



Trivial classifier: always predict majority class

Accuracy of a trivial classifier is: $16/20 = 80\%$

Random classifier: predict class 0 with probability 0.8 and class 1 with probability 0.2


Accuracy of the random classifier: 68%

$(0.8^2 + 0.2^2 = 0.68)$



Precision and recall

- Example: when a search engine returns 30 pages only 20 of which were relevant while failing to return 40 additional relevant pages, its precision is $20/30 = 2/3$ while its recall is $20/60 = 1/3$.



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$


How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$



ROC Curve

Predict positive if
 $p(y=1|x) > \text{threshold}$




Threshold = 0

Threshold = 1




ROC Curve example

e.g., diseased people, healthy people
blood protein levels normally distributed
Parameter that changes: threshold




TP	FP
FN	TN
1	1






ROC space





ROC space

A			B			C		
TP=63	FP=28	91	TP=77	FP=77	154	TP=24	FP=88	112
FN=37	TN=72	109	FN=23	TN=23	46	FN=76	TN=12	88
100	100	200	100	100	200	100	100	200
TPR = 0.63			TPR = 0.77			TPR = 0.24		
FPR = 0.28			FPR = 0.77			FPR = 0.88		
PPV = 0.69			PPV = 0.50			PPV = 0.21		
F1 = 0.66			F1 = 0.61			F1 = 0.22		
ACC = 0.68			ACC = 0.50			ACC = 0.18		





Area under the curve

- AUC or AUCROC gives the area under the ROC curve
- AUC is equal to the probability that a classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one
- Some issues in using AUC to compare classifiers (see “Measuring classifier performance: a coherent alternative to the area under the ROC curve”, Hand, JMLR, 2009)
 - can give unequal importance to a FPR or TPR for different classifiers